

## ✓ LSTM

LSTMs are a type of Recurrent Neural Network (RNN) that excel at **capturing long-term dependencies** in sequential data. They address the problem of vanishing and exploding gradients encountered in traditional RNNs by incorporating gates that regulate the flow of information.

## ✓ Core Concepts

1. **Memory Cell:** Holds information across time steps.
2. **Gates:** Control how information flows:
  - **Forget Gate:** Decides what information to forget.
  - **Input Gate:** Decides what new information to store.
  - **Output Gate:** Decides what information to output.

### Mathematical Representation

For each time step  $t$ :

1. Forget gate:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

2. Input gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

3. Update the cell state:

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

4. Output gate:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$

Here,  $\sigma$  is the sigmoid function,  $\odot$  is element-wise multiplication, and  $W, b$  are learnable parameters.

```
import torch
import torch.nn as nn
```

```
# Define the LSTM model
class LSTMModel(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers=1):
        super(LSTMModel, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers

        # LSTM layer
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)
        # Fully connected layer
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        # Initialize hidden and cell states
        h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
        c0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)

        # LSTM forward pass
        out, _ = self.lstm(x, (h0, c0))

        # Pass through the fully connected layer (only take last time step)
        out = self.fc(out[:, -1, :])
        return out
```

```
# Parameters
input_size = 10    # Number of features
hidden_size = 20   # Number of hidden units
output_size = 1    # Regression output
sequence_length = 5
batch_size = 16
```

```
# Initialize the model, loss, and optimizer
model = LSTMModel(input_size, hidden_size, output_size)
criterion = nn.MSELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

```
# Dummy data
x = torch.randn(batch_size, sequence_length, input_size) #input
```

```
y = torch.randn(batch_size, output_size)
```

```
# Training step
output = model(x)
loss = criterion(output, y)
optimizer.zero_grad()
loss.backward()
optimizer.step()

print("Training step completed. Loss:", loss.item())
```

🔄 Training step completed. Loss: 1.518990397453308

## ✓ Explanation of the Code

### 1. LSTM Layer:

- nn.LSTM handles the sequential nature of the input.

### 2. Hidden and Cell States:

- Initialized as zero tensors at the start of each sequence.

### 3. Fully Connected Layer:

- Converts the hidden state at the last time step to the desired output size.

### 4. Dummy Data:

- Random tensors simulate input and output for demonstration purposes.

This basic setup can be extended for tasks like time-series forecasting, text generation, or sequence classification.

Suppose we are working with time-series data:

- If you're forecasting stock prices:
  1. batch\_size: Number of different stock sequences (e.g., prices of 16 stocks).
  2. sequence\_length: Number of past days you consider (e.g., 5 days).
  3. input\_size: Features per day (e.g., [opening price, closing price, volume] → 3 features).
  4. x is the input. According to this code above, x is a tensor of random values with shape (16, 5, 10).

## ✓ How Input Flows Through the LSTM

### 1. Time Steps (Sequence):

- The LSTM processes the data one time step at a time. For a single sequence,  $x[:, t, :]$  represents the features at the  $t$ -th time step.

### 2. Features (Input Size):

- At each time step, the LSTM receives a vector of size `input_size`. For example, if `input_size=10`, there are 10 features per time step.

### 3. Batching:

- The LSTM processes multiple sequences simultaneously (batching). In this case, the model processes 16 sequences (one for each batch entry).

Input example- if,

- `batch_size = 2` # Two sequences in parallel
- `sequence_length = 4` # Four time steps per sequence
- `input_size = 3` # Three features per time step

Then  $x$  could look like-

- `x = torch.tensor([ [[1.0, 0.5, 2.0], [1.1, 0.4, 2.1], [1.2, 0.3, 2.2], [1.3, 0.2, 2.3]], [[0.9, 0.7, 1.8], [1.0, 0.6, 1.9], [1.1, 0.5, 2.0], [1.2, 0.4, 2.1]] ])`
- Shape: (2, 4, 3)

Here:

1. Sequence 1 has values `[[1.0, 0.5, 2.0], ...]` for its 4 time steps.
2. Sequence 2 has values `[[0.9, 0.7, 1.8], ...]` for its 4 time steps.

## Output of the LSTM

The LSTM processes the sequences and returns:

### 1. Output Tensor:

- A 3D tensor of shape (batch\_size, sequence\_length, hidden\_size) containing hidden states for each time step.

### 2. Final Hidden State:

- Used in this example to predict the output.

